

The education-obesity link across countries over time

Anastasios Goudras(2855565), Thedor(), Stefan(), Omar(), Ingmar(), Diego Melendez(2841772)

09 June 2026

Contents

Title Page	2
Part 1 - Identify a Social Problem	2
1.1 Describe the Social Problem	2
Part 2 - Data Sourcing	2
2.1 Load in the data	2
2.2 Provide a short summary of the dataset(s)	3
2.3 Describe the type of variables included	3
Part 3 - Quantifying	4
3.1 Data cleaning	4
Cleaning Step 1: Filter and clean the first dataset	4
Cleaning Step 2: Filter and clean the second dataset	4
Cleaning Step 3: Combine both datasets into one master file	5
3.2 Generate necessary variables	5
3.3 Visualize temporal variation	6
3.4 Visualize spatial variation	6
3.5 Visualize sub-population variation	6
3.6 Event analysis	6
Part 4 - Discussion	6
4.1 Discuss your findings	6
Part 5 - Reproducibility	6
5.1 Github repository link	6
5.2 Reference list	6

Title Page

Tutorial group: 6.1

Tutorial lecturer: Simon Loop

Part 1 - Identify a Social Problem

The social problem we have identified is the link between the BMI Index and educational attainment level, bridging two data sets from Eurostat to compare and identify any correlation between the two factors.

1.1 Describe the Social Problem

The relevance of these two variables, BMI and education levels, and how they are correlated with each other lies in the fact that obesity comes with serious health risks, such as diabetes and cardiovascular issues, which are often more common among people with lower educational attainment (World Health Organization, 2022; European Commission, 2024). This then results in the burden of health risks being disproportionately put on the lower educated class of society, and further increases social inequality issues. Beyond the health risks themselves, education has a clear impact on the ability of a person to maintain a healthy lifestyle, as higher education often comes with higher income, better access to higher quality food and so on. This means that addressing obesity cannot be done through increasing awareness alone, and that broader structural interventions, such as food pricing policies or workplace conditions need to be considered.

The fact we are looking at the European aspect of this issue also adds to its relevance. By looking at multiple countries over time, it becomes possible to assess whether this education gap in obesity rates is a universal pattern across Europe or whether certain countries have managed to narrow it. This kind of cross country comparison is useful for understanding what policy approaches may actually work in practice, especially since educational inequalities in obesity have been documented across Europe in recent research (World Health Organization, 2022; European Commission, 2024).

While the existence of an educational gradient in obesity is well established in the literature (World Health Organization, 2022; European Commission, 2024), far less is settled about how this gap behaves over time and whether it is narrowing or widening across different European countries. Much existing work documents the association at a single point in time or within individual countries, which leaves the question of long-term, cross country trends comparatively underexplored. This is the gap our project addresses: by tracking obesity rates by education level across multiple European countries and survey years, we aim to establish whether the education obesity divide is a stable, universal pattern or one that some countries have managed to reduce.

Part 2 - Data Sourcing

2.1 Load in the data

Both datasets are published by Eurostat. We downloaded them once and saved them as CSV files in the `data/` folder of our repository, which we read into R with `read_csv()`. Using saved copies rather than querying Eurostat live means the report always knits from the same data and does not depend on a network connection or on the Eurostat servers being reachable. The files are small, so storing them in the repository keeps the project self-contained and reproducible: anyone who clones it can knit the report without extra setup.

```
# Load the saved datasets from our repository
bmi_edu <- read_csv("data/bmi_edu.csv")
obesity_trend <- read_csv("data/obesity_trend.csv")
```

2.2 Provide a short summary of the dataset(s)

```
head(bmi_edu)
#> # A tibble: 6 x 9
#>   freq unit bmi isced11 sex age geo TIME_PERIOD values
#>   <chr> <chr> <chr> <chr> <chr> <chr> <chr> <dbl> <dbl>
#> 1 A PC BMI18P5-24 ED0-2 F TOTAL AT 2014 45.6
#> 2 A PC BMI18P5-24 ED0-2 F TOTAL AT 2019 40.9
#> 3 A PC BMI18P5-24 ED0-2 F TOTAL BE 2014 45.7
#> 4 A PC BMI18P5-24 ED0-2 F TOTAL BE 2019 41
#> 5 A PC BMI18P5-24 ED0-2 F TOTAL BG 2014 44.3
#> 6 A PC BMI18P5-24 ED0-2 F TOTAL BG 2019 48.1
head(obesity_trend)
#> # A tibble: 6 x 6
#>   freq unit bmi geo TIME_PERIOD values
#>   <chr> <chr> <chr> <chr> <dbl> <dbl>
#> 1 A PC BMI25-29 AL 2022 44.8
#> 2 A PC BMI25-29 AT 2008 36.5
#> 3 A PC BMI25-29 AT 2014 33.3
#> 4 A PC BMI25-29 AT 2017 35
#> 5 A PC BMI25-29 AT 2019 35
#> 6 A PC BMI25-29 AT 2022 35
```

Our first dataset, `bmi_edu` (Eurostat code `hlth_ehis_bm1e`, Body mass index (BMI) by sex, age and educational attainment level), comes from the European Health Interview Survey (EHIS), in which respondents report their own height and weight (Eurostat, n.d.-a). For each country and year it gives the share of people in each BMI category, broken down by education level, sex, and age. EHIS runs only every few years, so the data covers mainly the 2014 and 2019 waves. This is the dataset that links weight to education, so it is central to our analysis.

Our second dataset, `obesity_trend` (Eurostat code `sdg_02_10`, Obesity rate by body mass index (BMI)), is one of the EU's Sustainable Development Goal indicators, also based on self-reported health surveys (Eurostat, n.d.-b). It does not split the figures by education, but it covers many more years (2008, 2014, 2017, 2019, 2022), which makes it well suited to tracking obesity over time. It supplies the “across countries, over time” side of our research question. The data are reported by country, covering the EU member states together with several additional European countries, each shown as a two-letter code in the `geo` column.

2.3 Describe the type of variables included

Table 1: Variables in the two datasets.

Variable	Type	Dataset	Meaning
<code>freq</code>	character	both	data frequency (A = annual)
<code>unit</code>	character	both	unit of measurement (PC = percentage)
<code>bmi</code>	character	both	BMI category (e.g. overweight, obese)

Variable	Type	Dataset	Meaning
isced11	character	bmi_edu	education level (ISCED 2011 group)
sex	character	bmi_edu	sex (M, F, T = total)
age	character	bmi_edu	age group
geo	character	both	country (Eurostat code)
TIME_PERIOD	numeric	both	survey year
values	numeric	both	share of people in that category

As Table 1 shows, the variables fall into three groups. The health measure is BMI/obesity itself, the share of people in each weight category (the **bmi** and **values** columns). The socioeconomic measure is education (the **isced11** column in **bmi_edu**), which is the variable at the heart of our question. The remaining columns (**sex**, **age**, **geo** for country, and **TIME_PERIOD** for year) are background characteristics that we use to split the figures by group, place, and time.

Neither dataset comes from administrative records; both rely on surveys in which people report their own height and weight. This has two implications. First, self reported data tends to underestimate obesity slightly, since people often round their weight down. Second, the two datasets are not fully independent: the SDG indicator is partly built on the same EHIS survey as **bmi_edu**, so they are related rather than separate sources. We combined them deliberately: **bmi_edu** provides the education breakdown, while **obesity_trend** adds the longer time span. Because both are self reported survey data, our analysis can describe associations (for example, that obesity is more common among lower educated groups) but cannot show that education causes obesity.

Part 3 - Quantifying

3.1 Data cleaning

Stefan add explanation here

Cleaning Step 1: Filter and clean the first dataset

```
bmi_edu_clean <- bmi_edu %>%
  filter(
    bmi == "BMI_GE30", # Keep only Obese rows
    isced11 %in% c("ED0-2", "ED3-4", "ED5-8"), # Separate the 3 education groups
    sex == "T", # Combined sexes
    age == "TOTAL" # Combined adult age group
  ) %>%
  select(geo, TIME_PERIOD, isced11, bmi_value = values) %>%
  drop_na(bmi_value)
```

and here

Cleaning Step 2: Filter and clean the second dataset

```
obesity_trend_clean <- obesity_trend %>%
  filter(
    bmi == "BMI_GE30" # Keep only Obese rows since that's what we are comparing
  ) %>%
  select(geo, TIME_PERIOD, trend_value = values) %>%
  drop_na(trend_value)
```

and here

Cleaning Step 3: Combine both datasets into one master file

```
combined_dataset <- inner_join(bmi_edu_clean, obesity_trend_clean, by = c("geo", "TIME_PERIOD"))
```

and here

Please use a separate 'R block' of code for each type of cleaning. So, e.g. one for missing values, a new one for removing unnecessary variables etc.

3.2 Generate necessary variables

and here

Variable 1

```
# Create Variable 1: Human-readable education levels
bmi_edu_with_var1 <- combined_dataset %>%
  mutate(
    edu_level = case_when(
      isced11 == "ED0-2" ~ "Low Education",
      isced11 == "ED3-4" ~ "Medium Education",
      isced11 == "ED5-8" ~ "High Education"
    )
  )
```

```
# Reshaping before subtracting
gap_data <- bmi_edu_with_var1 %>%
  # 1. Isolate the columns we need
  select(geo, TIME_PERIOD, edu_level, bmi_value) %>%

  # 2. Flip the rows into columns
  pivot_wider(
    names_from = edu_level,
    values_from = bmi_value
  )
```

Variable 2

```
library(dplyr)

# Create Variable 2: Calculate the obesity inequality gap
final_gap_data <- gap_data %>%
  mutate(
    obesity_gap = `Low Education` - `High Education`
  )
```

3.3 Visualize temporal variation

3.4 Visualize spatial variation

Here you provide a description of why the plot above is relevant to your specific social problem.

3.5 Visualize sub-population variation

What is the obesity rate by country?

```
#{r visualise_subpopulations} dataset$inmetro <- dataset$inmetro %>% as.factor() #
Boxplot of poverty rate by state using the 'midwest' dataset ggplot(dataset, aes(x =
inmetro, y = percadultpoverty)) + geom_boxplot() + labs(title = "Distribution
of Poverty Rates by Metropolitan status (Midwest counties)", x = "Metropolitan Area",
y = "Poverty Rate of Adults (%)") + theme_minimal() + theme(legend.position =
"right")
```

Here you provide a description of why the plot above is relevant to your specific social problem.

3.6 Event analysis

Analyze the relationship between two variables (BMI .

Here you provide a description of why the plot above is relevant to your specific social problem.

Part 4 - Discussion

4.1 Discuss your findings

Part 5 - Reproducibility

5.1 Github repository link

<https://github.com/theodorkubi/ProgrammingForEcon>

5.2 Reference list

European Commission. (2024). Health inequalities: Overweight and obesity. Knowledge4Policy. https://knowledge4policy.ec.europa.eu/health-promotion-knowledge-gateway/health-inequalities-5_en

World Health Organization. (2022). Obesity and inequities. <https://iris.who.int/bitstreams/c025afdb-1ca7-49e2-ba89-adca1529e47c/download>

Eurostat. (n.d.-a). Body mass index (BMI) by sex, age and educational attainment level [Data set]. Retrieved June 7, 2026, from https://ec.europa.eu/eurostat/databrowser/view/hlth_ehis_bm1e

Eurostat. (n.d.-b). Obesity rate by body mass index (BMI) [Data set]. Retrieved June 7, 2026, from https://ec.europa.eu/eurostat/databrowser/view/sdg_02_10